

## Lecture 21 — March 6, 2019

*Patrick Lam**version 1*

Recall that we've been discussing beliefs. Here are a couple of beliefs that are worthwhile to check. (examples courtesy Dawson Engler.)

**Redundancy Checking.** 1) Code ought to do something. So, when you have code that doesn't do anything, that's suspicious. Look for identity operations, e.g.

$$x = x, 1 * y, x \& x, x | x.$$

Or, a longer example:

```
1      /* 2.4.5-ac8/net/appletalk/aarp.c */
2      da.s_node = sa.s_node;
3      da.s_net = da.s_net;
```

Also, look for unread writes:

```
1      for (entry=priv->lec_arp_tables[i];
2          entry != NULL; entry=next) {
3          next = entry->next; // never read!
4          ...
5      }
```

Redundancy suggests conceptual confusion.

So far, we've talked about MUST-beliefs; violations are clearly wrong (in some sense). Let's examine MAY beliefs next. For such beliefs, we need more evidence to convict the program.

**Process for verifying MAY beliefs.** We proceed as follows:

1. Record every successful MAY-belief check as “check”.
2. Record every unsuccessful belief check as “error”.
3. Rank errors based on “check” : “error” ratio.

Most likely errors occur when “check” is large, “error” small.

**Example.** One example of a belief is use-after-free:

```
1      free(p);
2      print(*p);
```

That particular case is a MUST-belief. However, other resources are freed by custom (undocumented) free functions. It's hard to get a list of what is a free function and what isn't. So, let's derive them behaviourally.

**Inferring beliefs: finding custom free functions.** The key idea is: if pointer `p` is not used after calling `foo(p)`, then derive a MAY belief that `foo(p)` frees `p`.

OK, so which functions are free functions? Well, just assume all functions free all arguments:

- emit “check” at every call site;
- emit “error” at every use.

(in reality, filter functions with suggestive names).

Putting that into practice, we might observe:

|                     |                     |                     |                     |                     |                     |
|---------------------|---------------------|---------------------|---------------------|---------------------|---------------------|
| <code>foo(p)</code> | <code>foo(p)</code> | <code>foo(p)</code> | <code>bar(p)</code> | <code>bar(p)</code> | <code>bar(p)</code> |
| <code>p = x;</code> | <code>p = x;</code> | <code>p = x;</code> | <code>p = 0;</code> | <code>p=0;</code>   | <code>p = x;</code> |

We would then rank `bar`’s error first. Plausible results might be: 23 free errors, 11 false positives.

**Inferring beliefs: finding routines that may return NULL.** The situation: we want to know which routines may return NULL. Can we use static analysis to find out?

- sadly, this is difficult to know statically (“`return p->next;`”?) and,
- we get false positives: some functions return NULL under special cases only.

Instead, let’s observe what the programmer does. Again, rank errors based on checks vs non-checks. As a first approximation, assume **all** functions can return NULL.

- if pointer checked before use: emit “check”;
- if pointer used before check: emit “error”.

This time, we might observe:

|                            |                              |                              |                              |
|----------------------------|------------------------------|------------------------------|------------------------------|
| <code>p = bar(...);</code> | <code>p = bar(...);</code>   | <code>p = bar(...);</code>   | <code>p = bar(...);</code>   |
| <code>p = x;</code>        | <code>if (!p) return;</code> | <code>if (!p) return;</code> | <code>if (!p) return;</code> |
|                            | <code>p = x;</code>          | <code>p = x;</code>          | <code>p = x;</code>          |

Again, sort errors based on the “check”:“error” ratio.

Plausible results: 152 free errors, 16 false positives.

## General statistical technique

When we write “a(); ... b();”, we mean a MAY-belief that a() is followed by b(). We don’t actually know that this is a valid belief. It’s a hypothesis, and we’ll try it out. Algorithm:

- assume every a–b is a valid pair;
- emit “check” for each path with “a()” and then “b()”;
- emit “error” for each path with “a()” and no “b()”.

(actually, prefilter functions that look paired).

Consider:

|                       |  |                       |  |                        |
|-----------------------|--|-----------------------|--|------------------------|
| foo(p, ...);          |  | foo(p, ...);          |  | foo(p, ...);           |
| bar(p, ...); // check |  | bar(p, ...); // check |  | // error: foo, no bar! |

This applies to the course project as well.

```
1 void scope1() {
2   A(); B(); C(); D();
3 }
4
5 void scope2() {
6   A(); C(); D();
7 }
8
9 void scope3() {
10  A(); B();
11 }
12
13 void scope4() {
14   B(); D(); scope1();
15 }
16
17 void scope5() {
18   B(); D(); A();
19 }
20
21 void scope6() {
22   B(); D();
23 }
```

“A() and B() must be paired”:  
either A() then B() or B() then A().

**Support** = # times a pair of functions appears together.  
 $\text{support}(\{A, B\}) = 3$

**Confidence**( $\{A, B\}, \{A\}$ ) =  
 $\text{support}(\{A, B\}) / \text{support}(\{A\}) = 3/4$

Sample output for support threshold 3, confidence threshold 65% (intra-procedural analysis):

- bug:A in scope2, pair: (A B), support: 3, confidence: 75.00%
- bug:A in scope3, pair: (A D), support: 3, confidence: 75.00%
- bug:B in scope3, pair: (B D), support: 4, confidence: 80.00%
- bug:D in scope2, pair: (B D), support: 4, confidence: 80.00%

The point is to find examples like the one from `cmpci.c` where there’s a `lock_kernel()` call, but, on an exceptional path, no `unlock_kernel()` call.

**Summary: Belief Analysis.** We don’t know what the right spec is. So, look for contradictions.

- MUST-beliefs: contradictions = errors!
- MAY-beliefs: pretend they’re MUST, rank by confidence.

(A key assumption behind this belief analysis technique: most of the code is correct.)

**Further references.** Dawson R. Engler, David Yu Chen, Seth Hallem, Andy Chou and Benjamin Chelf. “Bugs as Deviant Behaviors: A general approach to inferring errors in systems code”. In SOSP ’01.

Dawson R. Engler, Benjamin Chelf, Andy Chou, and Seth Hallem. “Checking system rules using system-specific, programmer-written compiler extensions”. In OSDI ’00 (best paper). [www.stanford.edu/~engler/mc-osdi.pdf](http://www.stanford.edu/~engler/mc-osdi.pdf)

Junfeng Yang, Can Sar and Dawson Engler. “eXplode: a Lightweight, General system for Finding Serious Storage System Errors”. In OSDI’06. [www.stanford.edu/~engler/explode-osdi06.pdf](http://www.stanford.edu/~engler/explode-osdi06.pdf)